

CSC 420 Project Report

Chen Wang and Jiaguan Tang

Utorid: wangc425 & tangjiag

Student #: 1006058926 & 1006488793

Table of Contents

1 Introduction -----	3
2 Shot Detection-----	3
2.1 Method-----	3
2.2 Implementation-----	3
2.3 Evaluation-----	4
3 Face Detection-----	6
3.1 Model choice-----	6
Haar Cascades Classifier-----	6
Dlib Recognizer-----	6
3.2 Capture Result and Compare:-----	7
4 Face matching-----	8
4.1 Methods choose-----	8
Harr Cascade Classifier-----	8
EigenFaceRecognizer-----	8
LBP Recognizer-----	8
Fisher Recognizer-----	9
Practice-----	9
Compare-----	9
4.2 Dlib Recognizer and face_recognition-----	9
4.3 Result-----	10
5 Contribution -----	11
6 Reference-----	12
7 Appendix-----	13

1. Introduction:

In this project, we are going to analyze movie trailers. Video analysis in computer vision focuses on object detection, which involves identifying and locating objects within a video frame. It has applications in surveillance, autonomous driving, and robotics. Various techniques, including deep learning-based methods, are used to improve the accuracy and effectiveness of object detection algorithms. In this scenario, we will be using shot detection and face detection techniques to analyze three different movie trailers.

The main logic of this project is divided into two parts. The first part is to go through the video and then get video clips. The second part is to track actor faces from the video shots and box the faces in the corresponding video clip.

2. Shot Detection

A shot refers to a series of consecutive video frames that are captured with smooth and continuous camera movement, and together form an uninterrupted sequence of visual content in terms of space, time, and graphical arrangement. The point at which one shot ends and another begins is known as a shot boundary. Shot detection is to detect shot boundaries within a video. It can help face detection by segmenting the video into different shots, allowing for the efficient application of face detection algorithms optimized for specific shot types and lighting conditions. [1] [2]

2.1 Method

There are several common shot detection methods.

Threshold-based methods: These methods analyze changes in visual content such as camera angle, lighting, and motion, and set a threshold for the change required to signal a new shot.

Histogram-based methods: These methods analyze the distribution of pixel values within the video frames and use changes in the distribution to identify shots.

Edge-based methods: These methods analyze the edges and contours of objects within the video frames and use changes in the edges to detect shots. [5]

2.2 Implementation

I choose Threshold-based methods. Choose a feature to use for shot detection. An example of a feature is frame differences.

2.2.1 absolute difference

One such method involves using the absolute difference between two consecutive frames as a feature to determine the shots by comparing them with a threshold value. However, determining a suitable threshold

value for this feature can be challenging. The difference between frames and different videos can vary significantly, which means that an independent threshold value must be set for each video. For instance, if video A is 1080p, and video B is 480p, the absolute difference in video A is much higher than that in video B. Despite using independent threshold values, the output can still be inaccurate since the changing pattern in one video can vary substantially. As such, it is essential to explore other features or methods that can provide more accurate shot detection.

2.2.2 average absolute difference

Another way is to use the average absolute difference between two consecutive frames as a feature. This approach is advantageous because it is independent of the video's resolution and can be used with various types of videos. However, this method is still prone to inaccuracies since the average difference can be high even when there is no shot boundary. For instance, in consecutive frames, the average absolute difference can be all high, like 16, or it can be very low like 2, which can result in false positives or negatives in shot detection. To overcome this limitation, we need to explore other features.

2.2.3 difference of average absolute difference

This feature calculates the difference of the average absolute difference between frames and uses the changing difference of frames to detect shot boundaries, aiming to avoid any biases from previous methods. A shot is defined as a continuous sequence of video frames captured by a single camera without any cuts or edits. Therefore, in a shot, the differences between all consecutive frames tend to be similar. For example, if the consecutive difference of four frames is 16, 15, and 16, then no shot boundary is detected. However, if the consecutive difference of four frames is 16, 70, and 70, a shot boundary is detected, and the previous shot ends in frame one, and the next shot starts in frame two. This feature has demonstrated a higher accuracy and outperforms previous methods.

2.3 evaluation

We will be using Recall, Precision, and F1 Score to determine the metrics of our shot detection. To calculate R, P, and F1 for shot detection, we need to define what constitutes a "positive" and "negative" case. In the context of shot detection, we can consider a positive case as a frame is a shot boundary and a negative case as a frame or time interval is not a shot boundary.

Manually compute a ground truth dataset that contains the actual shot boundaries, and a predicted dataset that contains the shot boundaries detected by our model

True Positives(Corrects): Frames that are correctly identified as containing a shot boundary.

False Positives (Falses): Frames that are incorrectly identified as containing a shot boundary.

False Negatives (Misses): Frames that are missed by the model and do contain a shot boundary.

True Negatives (TN): Frames that are correctly identified as not containing a shot boundary.

The mathematical expression of Recall, Precision, and F1 Score is:

$$Recall = \frac{Correct}{Correct + Miss}$$

$$Precision = \frac{Correct}{Correct + False}$$

$$F1 Score = 2 * \frac{Precision * Recall}{Precision + Recall}$$

Table1:

Trailer	C	M	F	R (%)	P (%)	F1 (%)
Trailer1	29	1	2	96.67	93.55	95.08
Trailer2	22	1	2	95.65	91.67	93.62
Trailer3	11	3	0	78.57	100	88.00

Figure 1

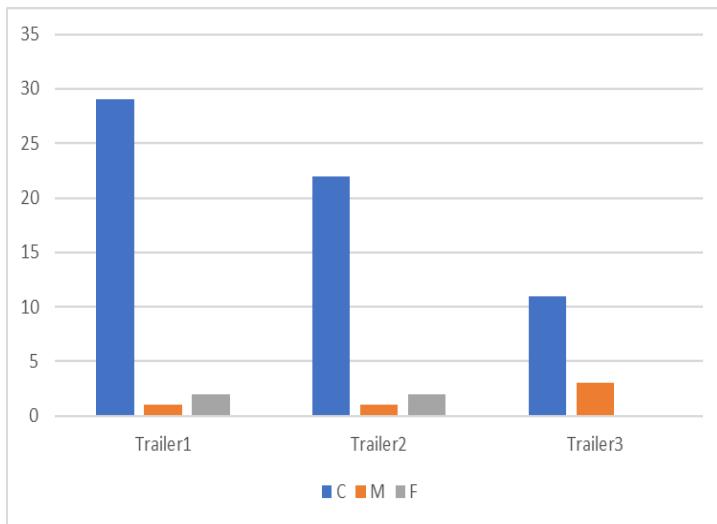
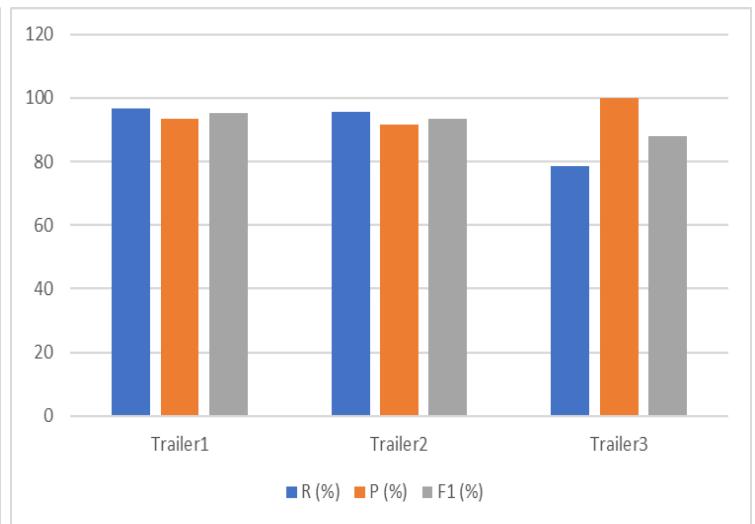


Figure 2



The table (**Table 1**) and graphs (**Figure1**, **Figure2**) show the metrics of our algorithm. As we can see the algorithm has a good metric overall. Most of the boundary can be detected if the camera break is cut. However, for the miss and false cases, I find that it is because some camera breaks of video are gradual transitions. During a gradual transition, the frame will not change significantly but gradually. The detector can't detect the boundary if the transition is smooth. As we can see from **Figure 3**.

Figure 3



3. Face detection

3.1 Model choice

There have been so many methods born for face detection. Some methods using deep learning concepts may have extraordinary performances compared to some older ideas. To compare the difference between them, I designed two models. The first one is called haar cascades which is first raised in 2001 and the second one is called dlib which is first raised in 2002.

3.1.1 Haar Cascades Classifier

The first model we are going to use is harr cascades. We know that some of the most common features on human faces are as follows:

Haar feature:

- *Eyes are darker than cheeks*
- *The bridge of the nose area is brighter compared to the eyes*
- *The positions of the eyes, mouth, and nose are relatively fixed*

For example, the first feature we will measure is the difference in intensity between the eye and the upper cheek. The method of eigenvalue calculation is very simple, just sum the pixels in the black area and subtract the pixels in the white area. And then, we apply this rectangle as a convolution kernel to the entire image. In order not to miss, we need to use all dimensions and positions of each convolution kernel.

To speed up this process, we use Adaboost. Once the features are extracted, apply them to the training set using an Adaboost classifier. And then, we use the haar cascade classifier as our face detector and it detects faces by the method “detectMultiScale”.

The following is a list of common parameters for the detectMultiScale function:

- *scaleFactor: Determines the scale size of each image.*
- *minNeighbors: Determines how many neighbouring boxes should be kept for each candidate rectangle.*
- *minSize: The size of the smallest object. Targets smaller than this value will be ignored.*
- *maxSize: The size of the largest object. Targets larger than this value will be ignored.*

Figure 4

```
faces = face_cascade.detectMultiScale(gray, 1.1, 5, minSize=(15,15))
```

3.1.2 Dlib Recognizer

The second model we used is Dlib. Dlib has two methods: hog and Cnn.

Hog vs Cnn: hog is faster, CNN is more accurate

We pass two parameters to the `face_recognition.face_locations` method:

RGB: our RGB image.

model: CNN or hog (the value is included in the command line arguments dictionary associated with the "detection_method" key). The CNN approach is more accurate but slower. HOG is faster but less accurate.

No need to tune and `face_recognition.face_location()` can directly locate the faces. [4]

Figure 5

```
boxes = face_recognition.face_locations(rgb)
encodings = face_recognition.face_encodings(rgb, boxes)
names = [1]
```

3.2 Capture Result and Compare:

In terms of speed, HOG is the fastest algorithm, followed by the Haar cascade classifier and CNN.

However, CNN in Dlib is the most accurate algorithm. HOG also performs well but has some issues recognizing smaller faces. The overall performance of Haar cascade classifiers is similar to HOG.

Below is the compared result: left is detected by Haar cascade and right is detected by Dlib recognizer

Figure 6

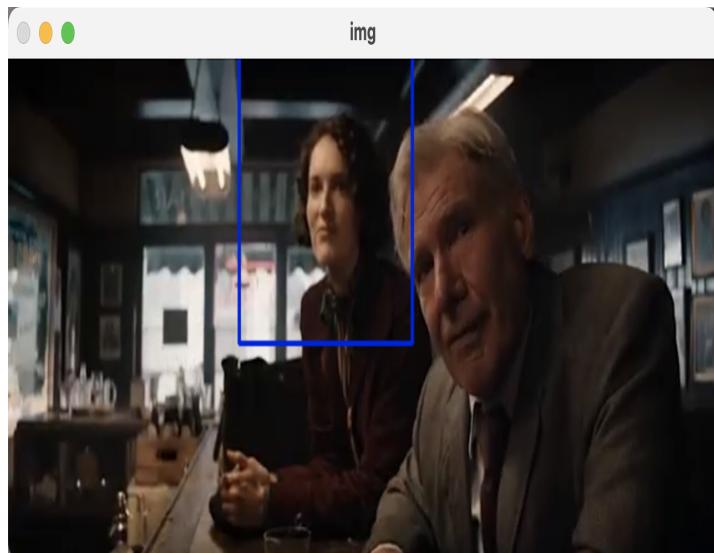
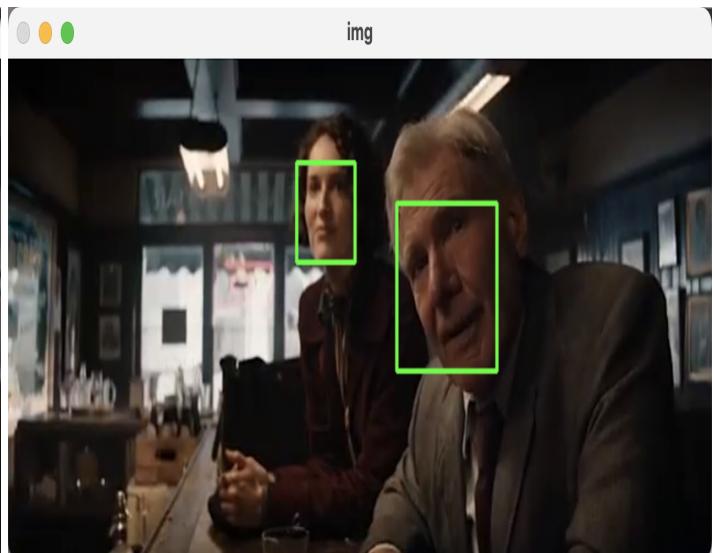


Figure 7



4. Face matching

4.1 Methods choose

4.1.1 Harr Cascade Classifier

For the Harr Cascade Classifier, we choose three recognizers to predict the actor faces in video shots: LBPHFaceRecognizer, EigenFaceRecognizer, and FisherFaceRecognizer.

We use `lbphFaceRecognizer.train()`, `EigenFaceRecognizer.train()`, and `FisherFaceRecognizer.train()` respectively to train the model to store the trained file as “model.yml” inside the trainer folder. [5]

Figure 8

```
lbph_recognizer = cv2.face.LBPHFaceRecognizer_create(1, 8, 8, 8, 123)
eigen_recognizer = cv2.face.EigenFaceRecognizer_create(10)
fisher_recognizer = cv2.face.FisherFaceRecognizer_create(10)
```

4.1.2 EigenFaceRecognizer:

First convert a batch of face images into a set of feature vectors, called "Eigenfaces", or "Eigenfaces", which contains the basic components of the initial training image set. The process of recognition is to project a new image into the eigenface subspace and pass its projection point in the subspace and the projection length of the line is used for judgment and identification.

After the image is transformed into another space, the images of the same category converge together, and the convergence of images of different categories is far away. In the original pixel space, it is difficult to divide the distribution of images of different categories with simple lines or planes. The specific implementation is to perform eigenvalue decomposition on the covariance matrix of all face images in the training set to obtain the corresponding eigenvectors, these eigenvectors are the "eigenfaces". Each eigenvector or eigenface is equivalent to capturing or describing a change or characteristic between human faces. This means that each face can be represented as a linear combination of these eigenfaces.

4.1.3 LBP Recognizer

The core idea of LBP is: to use the gray value of the central pixel as the threshold and compare it with its domain to obtain the corresponding binary code to represent the local texture features.

LBP extracts local features as the basis for discrimination. The significant advantage of the LBP method is that it is not sensitive to illumination, but it still does not solve the problem of pose and expression. However, compared with the eigenface method, the recognition rate of LBP has been greatly improved.

4.1.4 Fisher Recognizer

Because the set of feature vectors extracted by the fisher recognizer emphasizes the differences between different faces rather than changes in lighting conditions, facial expressions, and orientations, it is less sensitive to changes in illumination and facial posture for face recognition, which helps to improve the recognition effect.

4.1.5 Practice

After training each of the recognizers, we use them to predict a label based on the training set. Besides, setting a threshold to judge if the prediction is convincing or not. Like below:

Figure 9

```
if not face_test is None:  
    predictPCA, conf = eigen_recognizer.predict(face_test)  
    if conf > 90:  
        print("Predicted")
```

4.1.6 Compare

Based on the training data set, LBP and Eigen have a similar result while Fisher cannot recognize the actor/actress faces correctly.

Eigenface is not computationally efficient for large datasets because of SVD decomposition, the faces in the training pictures need to be aligned, and the pictures have the same size.

LBP cannot solve the problem of pose and expression.

4.2 Dlib Recognizer and face_recognition

We use a loop through the video and use the shot_detect function to detect if there is a face or not. If there is not, the loop will continue to the next frame. If there is a face detected, we use face_recognition.face_locations(image) to locate the face location and use encodings = face_recognition.face_encodings(RGB, boxes)to find similar faces from our dataset.

Figure 10

```
boxes = face_recognition.face_locations(rgb)  
encodings = face_recognition.face_encodings(rgb, boxes)
```

Here, we will use a dictionary to store the score and decide which actor/actress the face belongs to based on the score. After recognizing the face, we will use cv2.rectangle to draw a square outside the face and add its corresponding name to it. [3] [6]

Figure 11

```
matchedIdxs = [i for (i, b) in enumerate(matches) if b]
counts = {}
# loop over the matched indexes and maintain a count for
# each recognized face
for i in matchedIdxs:
    name = data["names"][i]
    counts[name] = counts.get(name, 0) + 1
# determine the recognized face with the largest number
# of votes (note: in the event of an unlikely tie Python
# will select first entry in the dictionary)
name = max(counts, key=counts.get)
```

4.3 Result

Use 45 photos per actor/actress to train the Haar Cascade classifier(9 actors/actresses in total) and the test accuracy is around 30% using LBP and Eigen recognizer. The accuracy is much lower(around 15%) using the fisher recognizer.

Use 14 photos per actor/actress to train the Dlib classifier(9 actors/actresses in total) and the test accuracy is around 90%.

This result agrees with the hypothesis we raised above: CNN in Dlib is the most accurate algorithm while Haar Cascade has a faster speed.

Below is the different result generated by Haar Cascade(left) and Face_recognition(right).

Left is wrong and right is correct. [7]

Figure 10

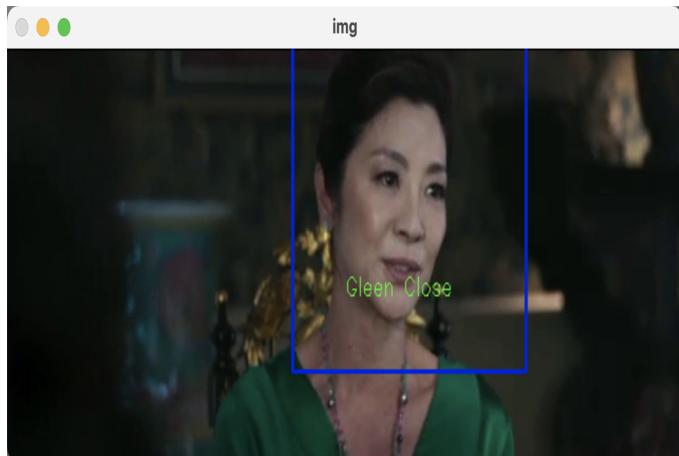


Figure 11

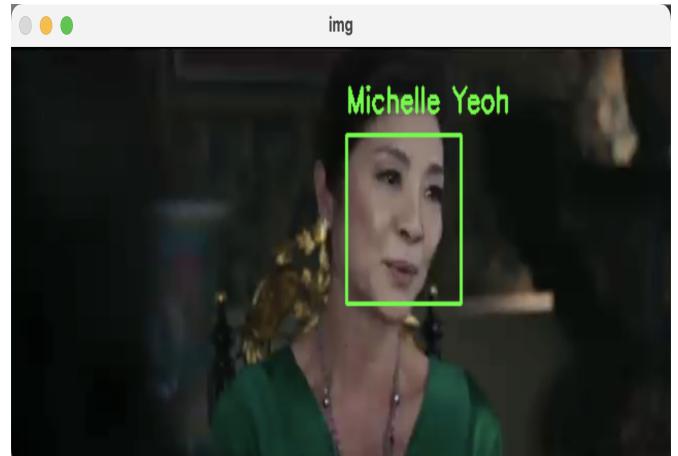
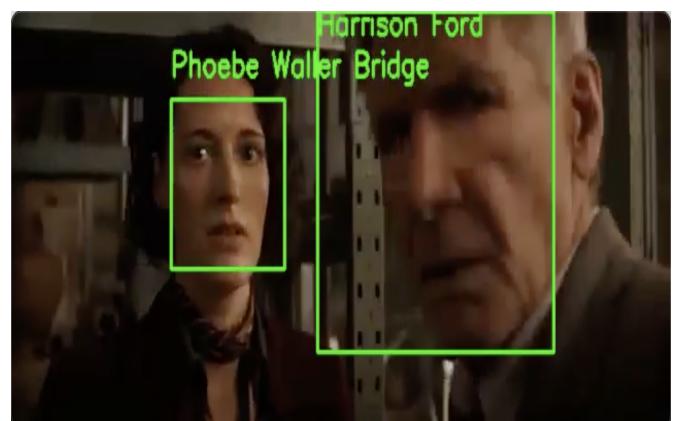


Figure 12



Figure 13



Contribution

Tang, Jiaguan:

- Shot Detection
- Data Collection

Wang, Chen:

- Face Detection
- Face Matching

Reference

1.

Zheng, G., & Xu, Y. (2021b). Efficient face detection and tracking in video sequences based on deep learning. *Information Sciences*, 568, 265–285.
<https://doi.org/10.1016/j.ins.2021.03.027>

2.

Ngo, C. (2009). Video Shot Detection. *Springer eBooks*, 3316–3320.

https://doi.org/10.1007/978-0-387-39940-9_1021

3.

Davisking. (n.d.). *GitHub - davisking/dlib-models: Trained model files for dlib example programs*. GitHub.
<https://github.com/davisking/dlib-models>

4.

Rosebrock, A. (2023, March 23). *Face detection with OpenCV and deep learning - PyImageSearch*.

PyImageSearch. <https://pyimagesearch.com/2018/02/26/face-detection-with-opencv-and-deep-learning/>

5.

Video Shot Detection Techniques Brief Overview

<http://www.ijergs.org/files/documents/VIDEO-108.pdf>

6.

Rosebrock, A. (2021, April 17). *Faster video file FPS with cv2.VideoCapture and OpenCV - PyImageSearch*.

PyImageSearch.

<https://pyimagesearch.com/2017/02/06/faster-video-file-fps-with-cv2-videocapture-and-opencv/>

7.

Suárez, I. (2021, December 27). Improving your image matching results by 14% with one line of code.

Medium.

<https://towardsdatascience.com/improving-your-image-matching-results-by-14-with-one-line-of-code-b72ae9ca2b73>

Appendix

Shot Detection Python code:

```
import cv2
from google.colab.patches import cv2_imshow
import matplotlib.pyplot as plt
import numpy as np

movie1 = cv2.VideoCapture("Movie_1_new.mp4")
movie2 = cv2.VideoCapture("Movie_2_new.mp4")
movie3 = cv2.VideoCapture("Movie_3_new.mp4")

def shot_detector(video):
    # Set the threshold for shot detection
    threshold = 7
    # Initialize variables
    frame_count = 0
    prev_frame = None
    shots = []
    dif = []
    start = True
    skip = True

    while True:
        # Read the next frame
        ret, frame = video.read()

        # If there are no more frames, break out of the loop
        if not ret:
            if shots != []:
                shots[-1][-1] = frame_count
            break

        # Convert the frame to grayscale
        gray_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)

        # If this is not the first frame, compare it to the previous frame
        if prev_frame is not None:

            diff = cv2.absdiff(gray_frame, prev_frame)

            if np.average(diff) > 0.2 or not skip:
                if len(dif) > 0:
                    if np.average(diff) - dif[-1] > threshold:
                        shots[-1][-1] = frame_count - 1
                        shots.append([frame_count, frame_count])
                dif.append(np.average(diff))
                skip = True
            elif skip:
                shots[-1][-1] = frame_count
                skip = False

        else:
            shots.append([frame_count, frame_count])
            start = False

        # Update variables
        frame_count += 1
        prev_frame = gray_frame

    return shots
```

```

def accu(output, truth):
    out_put = []
    truth_ = []
    correct = 0
    miss = 0
    false = 0
    for i in range(1, len(output)):
        out_put.append(output[i][0])
    for i in range(1, len(truth)):
        truth_.append(truth[i][0])
    for i in truth_:
        if i in out_put:
            correct += 1
        else:
            miss += 1
    for i in out_put:
        if i not in truth_:
            false += 1
    return correct, miss, false

o1 = shot_detector(movie1)
o2 = shot_detector(movie2)
o3 = shot_detector(movie3)

print("Video1_output: \n", o1)
print("Video2_output: \n", o2)
print("Video2_output: \n", o3)

Video1_output:
[[0, 50], [51, 92], [93, 123], [124, 192], [193, 222], [223, 261], [262, 300], [301, 316], [317, 339], [340, 391], [392, 412], [413, 429], [430, 491], [51, 77], [78, 109], [110, 149], [150, 211], [212, 250], [251, 278], [279, 296], [297, 328], [329, 360], [361, 391], [392, 418], [419, 446], [45, 80], [81, 156], [157, 191], [192, 233], [234, 274], [275, 315], [316, 345], [346, 492], [493, 522], [523, 586]]
Video2_output:
[[0, 51], [52, 77], [78, 109], [110, 149], [150, 211], [212, 250], [251, 278], [279, 296], [297, 328], [329, 360], [361, 391], [392, 418], [419, 446], [45, 80], [81, 156], [157, 191], [192, 233], [234, 274], [275, 315], [316, 345], [346, 381], [382, 424], [425, 469], [470, 492], [493, 522], [523, 586]]
Video2_output:
[[0, 4], [5, 45], [46, 80], [81, 156], [157, 191], [192, 233], [234, 274], [275, 315], [316, 345], [346, 492], [493, 522], [523, 586]]

t1 = [[0, 50], [51, 92], [93, 123], [124, 192], [193, 222], [223, 261], [262, 300], [301, 316], [317, 339], [340, 391], [392, 412], [413, 429], [430, 491], [51, 77], [78, 109], [110, 149], [150, 211], [212, 250], [251, 278], [279, 296], [297, 328], [329, 360], [361, 391], [392, 418], [419, 446], [45, 80], [81, 156], [157, 191], [192, 233], [234, 274], [275, 315], [316, 345], [346, 492], [493, 522], [523, 586]]
t2 = [[0, 51], [52, 77], [78, 109], [110, 149], [150, 211], [212, 250], [251, 278], [279, 296], [297, 328], [329, 360], [361, 391], [392, 418], [419, 446], [45, 80], [81, 156], [157, 191], [192, 233], [234, 274], [275, 315], [316, 345], [346, 381], [382, 424], [425, 469], [470, 492], [493, 522], [523, 586]]
t3 = [[0, 4], [5, 45], [46, 80], [81, 156], [157, 191], [192, 233], [234, 274], [275, 315], [316, 345], [346, 381], [382, 424], [425, 469], [470, 492], [493, 522], [523, 586]]

Movie1_accuracy = accu(o1, t1)
Movie2_accuracy = accu(o2, t2)
Movie3_accuracy = accu(o3, t3)

print("The num of correct, miss, false of Movie1: \n", Movie1_accuracy)
print("The num of correct, miss, false of Movie2: \n", Movie2_accuracy)
print("The num of correct, miss, false of Movie3: \n", Movie3_accuracy)

The num of correct, miss, false of Movie1:
(29, 1, 2)
The num of correct, miss, false of Movie2:
(22, 1, 2)
The num of correct, miss, false of Movie3:
(11, 3, 0)

```

Face Detection Python code:

```
1  import cv2
2  import numpy as np
3  from imutils.video import VideoStream
4  import face_recognition
5  import argparse
6  import imutils
7  import pickle
8  from imutils import paths
9  import argparse
10 import os
11 from PIL import Image
12
13 class dlib:
14     def __init__(self):
15         self.video_frames = []
16
17     def create(self):
18         ap = argparse.ArgumentParser()
19         ap.add_argument("-i", "--dataset", required=True,
20                         help="path to input directory of faces + images")
21         ap.add_argument("-e", "--encodings", required=True,
22                         help="path to serialized db of facial encodings")
23         ap.add_argument("-d", "--detection-method", type=str, default="cnn",
24                         help="face detection model to use: either `hog` or `cnn`")
25         args = vars(ap.parse_args())
26
27         # grab the paths to the input images in our dataset
28         print("[INFO] quantifying faces...")
29         imagePaths = list(paths.list_images(args["dataset"]))
30         # initialize the list of known encodings and known names
31
32         knownEncodings = []
33         knownNames = []
34
35         # loop over the image paths
36         for (i, imagePath) in enumerate(imagePaths):
37             # extract the person name from the image path
38             print("[INFO] processing image {}/{}".format(i + 1,
39                 len(imagePaths)))
40             name = 'unknown'
41             tag = int(imagePath.split(os.path.sep)[-1].split('.')[1])
42             if tag == 1:
43                 name = "Thomas Kretschman"
44             elif tag == 2:
45                 name = "Harrison Ford"
46             elif tag == 3:
47                 name = "Phoebe Waller Bridge"
48             elif tag == 4:
49                 name = "Michelle Yeoh"
50             elif tag == 5:
51                 name = "Henry Golding"
52             elif tag == 6:
53                 name = "Constance Wu"
54             elif tag == 7:
55                 name = "Jonathan Pryce"
56             elif tag == 8:
57                 name = "Max Irons"
58             elif tag == 9:
59                 name = "Gleen Close"
60             # name = imagePath.split(os.path.sep)[-2]
61             # load the input image and convert it from BGR (OpenCV ordering)
```

```

61     # to dlib ordering (RGB)
62     image = cv2.imread(imagePath)
63     rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
64
65     # detect the (x, y)-coordinates of the bounding boxes
66     # corresponding to each face in the input image
67     boxes = face_recognition.face_locations(rgb,
68         model=args["detection_method"])
69     # compute the facial embedding for the face
70     encodings = face_recognition.face_encodings(rgb, boxes)
71     # loop over the encodings
72     for encoding in encodings:
73         # add each encoding + name to our set of known names and
74         # encodings
75         knownEncodings.append(encoding)
76         knownNames.append(name)
77
78     # dump the facial encodings + names to disk
79     print("[INFO] serializing encodings...")
80     data = {"encodings": knownEncodings, "names": knownNames}
81     f = open(args["encodings"], "wb")
82     f.write(pickle.dumps(data))
83     f.close()
84
85 def phototest(self):
86     ap = argparse.ArgumentParser()
87     ap.add_argument("-e", "--encodings", required=True,
88                     help="path to serialized db of facial encodings")
89     ap.add_argument("-i", "--image", required=True,
90                     help="path to input image")

```

```

91     ap.add_argument("-d", "--detection-method", type=str, default="cnn",
92     help="face detection model to use: either `hog` or `cnn`")
93     args = vars(ap.parse_args())
94     # load the known faces and embeddings
95     print("[INFO] loading encodings...")
96     data = pickle.loads(open(args["encodings"], "rb").read())
97     # load the input image and convert it from BGR to RGB
98     image = cv2.imread(args["image"])
99     rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
100    # detect the (x, y)-coordinates of the bounding boxes corresponding
101    # to each face in the input image, then compute the facial embeddings
102    # for each face
103    boxes = face_recognition.face_locations(rgb,
104        model=args["detection_method"])
105    encodings = face_recognition.face_encodings(rgb, boxes)
106    # initialize the list of names for each face detected
107    names = []
108    # loop over the facial embeddings
109    for encoding in encodings:
110        # attempt to match each face in the input image to our known
111        # encodings
112        matches = face_recognition.compare_faces(data["encodings"],
113            encoding)
114        name = "Unknown"
115        # check to see if we have found a match
116        if True in matches:
117            # find the indexes of all matched faces then initialize a
118            # dictionary to count the total number of times each face
119            # was matched
120            matchedIdxs = [i for (i, b) in enumerate(matches) if b]
121
122            counts = {}
123            # loop over the matched indexes and maintain a count for
124            # each recognized face face
125            for i in matchedIdxs:
126                name = data["names"][i]
127                counts[name] = counts.get(name, 0) + 1
128            # determine the recognized face with the largest number of
129            # votes (note: in the event of an unlikely tie Python will
130            # select first entry in the dictionary)
131            name = max(counts, key=counts.get)
132
133            # update the list of names
134            names.append(name)
135            # loop over the recognized faces
136            for (top, right, bottom, left), name in zip(boxes, names):
137                # draw the predicted face name on the image
138                cv2.rectangle(image, (left, top), (right, bottom), (0, 255, 0), 2)
139                y = top - 15 if top - 15 > 15 else top + 15
140                cv2.putText(image, name, (left, y), cv2.FONT_HERSHEY_SIMPLEX,
141                    0.75, (0, 255, 0), 2)
142            # show the output image
143            cv2.imshow("Image", image)
144            cv2.waitKey(0)
145
146            def shot_detector(self, video_name):
147                video = cv2.VideoCapture(video_name)
148
149                # Set the threshold for shot detection
150                threshold = 6

```

```

151     # Initialize variables
152     frame_count = 0
153     prev_frame = None
154     shots = []
155     dif = []
156     start = True
157
158     while True:
159         # Read the next frame
160         ret, frame = video.read()
161
162         # If there are no more frames, break out of the loop
163         if not ret:
164             break
165
166         # Convert the frame to grayscale
167         gray_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
168
169         # If this is not the first frame, compare it to the previous frame
170         # print(prev_frame)
171         if prev_frame is not None:
172
173             diff = cv2.absdiff(gray_frame, prev_frame)
174
175             if len(dif) > 0:
176                 # print(np.average(diff) - dif[-1] > threshold,start)
177                 if dif[-1] - np.average(diff) > threshold and start:
178
179                     shots.append([frame_count, frame_count])
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210

```

```

210         frame_count += 1
211         prev_frame = gray_frame
212
213     # Print the shot boundaries
214     return shots
215
216 def has_face(self, img):
217     rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
218     rgb = imutils.resize(img, width=750)
219     boxes = face_recognition.face_locations(rgb)
220     if boxes:
221         return True
222     return False
223
224 def inShot(self, index, l):
225     for item in l:
226         if item[0] <= index <= item[1] and [self.has_face(self.video_frames[item[0]]), or
227                                         self.has_face(self.video_frames[item[1]])]:
228             return True
229     return False
230
231 def videotest(self):
232     ap = argparse.ArgumentParser()
233     ap.add_argument("-e", "--encodings", required=True,
234     help="path to serialized db of facial encodings")
235     ap.add_argument("-v", "--video", required=True,
236     help="path to input video file")
237     ap.add_argument("-o", "--output", type=str,
238     help="path to output video")
239     # ap.add_argument("-y", "--display", type=int, default=1,
240     # help="whether or not to display output frame to screen")
241
242     # ap.add_argument("-d", "--detection-method", type=str, default="cnn",
243     # help="face detection model to use: either `hog` or `cnn`")
244     args = vars(ap.parse_args())
245     # load the known faces and embeddings
246     data = pickle.loads(open(args["encodings"], "rb").read())
247     # initialize the video stream and pointer to output video file, then
248     # allow the camera sensor to warm up
249     vs = cv2.VideoCapture(args["video"])
250     # vs = VideoStream(src=0).start()
251     writer = None
252     video = cv2.VideoCapture(args["video"])
253     while True:
254         # Read the next frame
255         ret, frame = video.read()
256         if frame is not None:
257             self.video_frames.append(frame)
258             # If there are no more frames, break out of the loop
259             if not ret:
260                 break
261     video.release()
262     shotBoundaryIndex = self.shot_detector(args["video"])
263     i = 0
264     # loop over frames from the video file stream
265     while True:
266         # grab the frame from the threaded video stream
267         ret, frame = vs.read()
268         if not ret:
269             break
270         # convert the input frame from BGR to RGB then resize it to have
271         # a width of 750px (to speedup processing)

```

```

271     if self.inShot(i, shotBoundaryIndex):
272         rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
273         rgb = imutils.resize(frame, width=750)
274         r = frame.shape[1] / float(rgb.shape[1])
275         # detect the (x, y)-coordinates of the bounding boxes
276         # corresponding to each face in the input frame, then compute
277         # the facial embeddings for each face
278         boxes = face_recognition.face_locations(rgb)
279         encodings = face_recognition.face_encodings(rgb, boxes)
280         names = []
281         # loop over the facial embeddings
282         for encoding in encodings:
283             # attempt to match each face in the input image to our known
284             # encodings
285             matches = face_recognition.compare_faces(data["encodings"],
286                 encoding)
287             name = "Unknown"
288             # check to see if we have found a match
289             if True in matches:
290                 # find the indexes of all matched faces then initialize a
291                 # dictionary to count the total number of times each face
292                 # was matched
293                 matchedIdxs = [i for (i, b) in enumerate(matches) if b]
294                 counts = {}
295                 # loop over the matched indexes and maintain a count for
296                 # each recognized face face
297                 for i in matchedIdxs:
298                     name = data["names"][i]
299                     counts[name] = counts.get(name, 0) + 1
300                     # determine the recognized face with the largest number

```

```

301             # of votes (note: in the event of an unlikely tie Python
302             # will select first entry in the dictionary)
303             name = max(counts, key=counts.get)
304             if counts[name] <= 3 :
305                 name = 'Unknown'
306
307                 # update the list of names
308                 names.append(name)
309             # Loop over the recognized faces
310             for ((top, right, bottom, left), name) in zip(boxes, names):
311                 # rescale the face coordinates
312                 top = int(top * r)
313                 right = int(right * r)
314                 bottom = int(bottom * r)
315                 left = int(left * r)
316                 # draw the predicted face name on the image
317                 cv2.rectangle(frame, (left, top), (right, bottom),
318                               (0, 255, 0), 2)
319                 y = top - 15 if top - 15 > 15 else top + 15
320                 cv2.putText(frame, name, (left, y), cv2.FONT_HERSHEY_SIMPLEX,
321                             0.75, (0, 255, 0), 2)
322             i += 1
323             cv2.imshow('img',frame)
324             key = cv2.waitKey(1)
325             if key == ord('q'):
326                 break
327             # if the video writer is None *AND* we are supposed to write
328             # the output video to disk initialize the writer
329             if writer is None and args["output"] is not None:
330                 fourcc = cv2.VideoWriter_fourcc(*"MJPG")

```

```

331         writer = cv2.VideoWriter(args["output"], fourcc, 20,
332             (frame.shape[1], frame.shape[0]), True)
333         # if the writer is not None, write the frame with recognized
334         # faces to disk
335         if writer is not None:
336             writer.write(frame)
337         # check to see if we are supposed to display the output frame to
338         # the screen
339         # if args["display"] > 0:
340         #     cv2.imshow("Frame", frame)
341         #     key = cv2.waitKey(0) & 0xFF
342         #     # if the `q` key was pressed, break from the loop
343         #     if key == ord("q"):
344         #         break
345         # do a bit of cleanup
346         cv2.destroyAllWindows()
347         vs.release()
348         # check to see if the video writer point needs to be released
349         if writer is not None:
350             writer.release()
351
352
353 class haarcascade:
354     def __init__(self):
355         self.detector = cv2.CascadeClassifier(cv2.data.haarcascades + 'haarcascade_frontalface_alt.xml')
356         self.lbph_recognizer = cv2.face.LBPHFaceRecognizer_create(1, 8, 8, 8, 123)
357         self.eigen_recognizer = cv2.face.EigenFaceRecognizer_create(10)
358         self.fisher_recognizer = cv2.face.FisherFaceRecognizer_create(10)
359         self.videoName = 'video/Movie_2.mp4'
360         self.cascade_path = cv2.data.haarcascades + "haarcascade_frontalface_alt.xml"
361
362         self.face_cascade = cv2.CascadeClassifier(self.cascade_path)
363         self.video_frames = []
364
365     def process(self):
366         sampleNum = 1
367         Id = 0
368         pic = 1
369
370         while True:
371             if sampleNum > 14:
372                 sampleNum -= 14
373             if pic <= 14:
374                 Id = 1
375             elif 14 < pic <= 28:
376                 Id = 2
377             elif 28 < pic <= 42:
378                 Id = 3
379             elif 42 < pic <= 56:
380                 Id = 4
381             elif 56 < pic <= 70:
382                 Id = 5
383             elif 70 < pic <= 84:
384                 Id = 6
385             elif 84 < pic <= 98:
386                 Id = 7
387             elif 98 < pic <= 112:
388                 Id = 8
389             elif 112 < pic <= 126:
390                 Id = 9
391             if Id == 1:

```

```

391     filename = './actor/thomas kretschmann.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
392 elif Id == 2:
393     filename = './actor/harrison ford.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
394 elif Id == 3:
395     filename = './actor/phoebe waller bridge.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
396 elif Id == 4:
397     filename = './actor/Michelle Yeoh.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
398 elif Id == 5:
399     filename = './actor/Henry Golding.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
400 elif Id == 6:
401     filename = './actor/COnstance Wu.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
402 elif Id == 7:
403     filename = './actor/Jonathan Pryce.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
404 elif Id == 8:
405     filename = './actor/Max Irons.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
406 elif Id == 9:
407     filename = './actor/Glenn Close.' + str(Id) + '.' + str(sampleNum) + '.jpeg'
408 img = cv2.imread(filename)
409 gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
410 faces = self.detector.detectMultiScale(gray, 1.3, 5, minSize=(50,50))
411 if type(faces) == np.ndarray:
412     for (x, y, w, h) in faces:
413         cv2.rectangle(img, (x, y), (x + w, y + h), (255, 0, 0), 2)
414         f = cv2.resize(gray[y:y+h, x:x+w], (200,200))
415         # saving the captured face in the dataset folder
416         cv2.imwrite("dataSet/Actor." + str(Id) + '.' + str(sampleNum) + ".jpg", f)
417
418 cv2.imshow('frame', img)
419 # wait for 100 miliseconds
420 sampleNum += 1

```

```

421     pic += 1
422     if cv2.waitKey(50) & 0xFF == ord('q'):
423         break
424     if pic > 135:
425         break
426
427     cv2.destroyAllWindows()
428
429 def shot_detector(self, video_name):
430     video = cv2.VideoCapture(video_name)
431
432     # Set the threshold for shot detection
433     threshold = 6
434
435     # Initialize variables
436     frame_count = 0
437     prev_frame = None
438     shots = []
439     dif = []
440     start = True
441
442     while True:
443         # Read the next frame
444         ret, frame = video.read()
445
446         # If there are no more frames, break out of the loop
447         if not ret:
448             break
449
450         # Convert the frame to grayscale
451
452         gray_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
453
454         # If this is not the first frame, compare it to the previous frame
455         # print(prev_frame)
456         if prev_frame is not None:
457
458             diff = cv2.absdiff(gray_frame, prev_frame)
459
460             if len(dif) > 0:
461                 # print(np.average(diff) - dif[-1] > threshold,start)
462                 if dif[-1] - np.average(diff) > threshold and start:
463
464                     shots.append([frame_count, frame_count])
465
466                     # plt.imshow(prev_frame, cmap='gray')
467                     # plt.axis("off")
468                     # plt.title("start at " + str(frame_count))
469                     # plt.show()
470                     start = False
471                 elif np.average(diff) - dif[-1] > threshold and not start:
472
473                     # plt.imshow(prev_frame, cmap='gray')
474                     # plt.axis("off")
475                     # plt.title("end at " + str(frame_count))
476                     # plt.show()
477
478                     shots[-1][-1] = frame_count
479                     start = True
480                 dif.append(np.average(diff))

```

```

481
482     else:
483         shots.append([frame_count, frame_count])
484
485         # plt.imshow(gray_frame, cmap='gray')
486         # plt.axis("off")
487         # plt.title("start at " + str(frame_count))
488         # plt.show()
489
490         start = False
491         dif.append(0)
492
493         # Update variables
494         frame_count += 1
495         prev_frame = gray_frame
496
497         # Print the shot boundaries
498         return shots
499
500     def train(self):
501         faces, Ids = self.get_images_and_labels('dataSet')
502         self.lbph_recognizer.train(faces, np.array(Ids))
503         self.eigen_recognizer.train(faces, np.array(Ids))
504         self.fisher_recognizer.train(faces, np.array(Ids))
505         self.lbph_recognizer.save('haarcascades/trainner/lbph_trainner.yml')
506         self.eigen_recognizer.save('haarcascades/trainner/eigen_trainner.yml')
507         self.fisher_recognizer.save('haarcascades/trainner/fisher_trainner.yml')
508
509     def get_images_and_labels(self, path):
510         image_paths = [os.path.join(path, f) for f in os.listdir(path)]

```

```

511     face_samples = []
512     ids = []
513
514     for image_path in image_paths:
515         if image_path == 'dataSet/' + '.DS_Store':
516             continue
517         image = Image.open(image_path).convert('L')
518         image_np = np.array(image, 'uint8')
519         if os.path.split(image_path)[-1].split(".")[-1] != 'jpg':
520             continue
521         image_id = int(os.path.split(image_path)[-1].split(".")[1])
522         faces = self.detector.detectMultiScale(image_np, 1.1, 5, minSize=(50,50))
523         for (x, y, w, h) in faces:
524             face_samples.append(cv2.resize(image_np[y:y + h, x:x + w], (200,200)))
525             ids.append(image_id)
526
527     return face_samples, ids
528
529 def createframe(self):
530     video = cv2.VideoCapture(self.videoName)
531     while True:
532         # Read the next frame
533         ret, frame = video.read()
534         if frame is not None:
535             self.video_frames.append(frame)
536         # If there are no more frames, break out of the loop
537         if not ret:
538             break
539     video.release()
540
541 def has_face(self, img):
542     gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
543     faces = self.face_cascade.detectMultiScale(gray, 1.1, 5, minSize=(15,15))
544     if type(faces) == np.ndarray:
545         return True
546     return False
547
548 def inShot(self, index, l):
549     for item in l:
550         if item[0] <= index <= item[1] and [self.has_face(self.video_frames[item[0]]), or
551                                         | self.has_face(self.video_frames[item[1]])]:
552             return True
553     return False
554
555 def testshot(self):
556     self.lbph_recognizer.read('haarcascades/trainner/lbph_trainner.yml')
557     self.eigen_recognizer.read('haarcascades/trainner/eigen_trainner.yml')
558     self.fisher_recognizer.read('haarcascades/trainner/fisher_trainner.yml')
559
560     shotBoundaryIndex = self.shot_detector(self.videoName)
561     # font = cv2.cv.InitFont(cv2.cv.CV_FONT_HERSHEY_SIMPLEX, 1, 1, 0, 1, 1) # in OpenCV 2
562     font = cv2.FONT_HERSHEY_SIMPLEX
563     video = cv2.VideoCapture(self.videoName)
564     v = cv2.VideoCapture(self.videoName)
565     while True:
566         # Read the next frame
567         ret, frame = v.read()
568         if frame is not None:
569             self.video_frames.append(frame)
570         # If there are no more frames, break out of the loop

```

```

571     if not ret:
572         break
573     v.release()
574     i = 0
575     while True:
576         ret, img = video.read()
577         if not ret:
578             break
579         if self.inShot(i, shotBoundaryIndex):
580             gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
581             faces = self.face_cascade.detectMultiScale(gray, 1.2, 5, minSize=(15,15))
582             for (x, y, w, h) in faces:
583                 face = gray[y:y+h, x:x+w]
584                 # k1 = np.array([[y, x], [y+h,x], [y,x+w], [y+h,x+w]])
585                 # k2 = np.array([[0,0],[0,200],[200,0],[200,200]])
586                 # h,mark = cv2.findHomography(k1,k2,cv2.RANSAC)
587                 # face_test = cv2.warpPerspective(face, h, (200,200))
588                 # face_test = cv2.resize(face, (200, 200))
589                 cv2.rectangle(img, (x - 50, y - 50), (x + w + 50, y + h + 50), (225, 0, 0), 2)
590                 # img_id, conf = fisher_recognizer.predict(face_test)
591                 # if conf > 80:
592                     # if img_id == 1:
593                         # img_id = 'Thomas Kretschman'
594                     # elif img_id == 2:
595                         # img_id = 'Harrison Ford'
596                     # elif img_id == 3:
597                         # img_id = 'Phoebe Waller Bridge'
598                     # elif img_id == 4:
599                         # img_id = 'Shaunette Renée Wilson'
600                     # else:
601                         # img_id = "Unknown"
602                     predictPCA = 0
603                     if not face_test is None:
604                         predictPCA, conf = self.eigen_recognizer.predict(face_test)
605                     if conf > 90:
606                         if predictPCA == 1:
607                             name = "Thomas Kretschman"
608                         elif predictPCA == 2:
609                             name = "Harrison Ford"
610                         elif predictPCA == 3:
611                             name = "Phoebe Waller Bridge"
612                         elif predictPCA == 4:
613                             name = "Michelle Yeoh"
614                         elif predictPCA == 5:
615                             name = "Henry Golding"
616                         elif predictPCA == 6:
617                             name = "Constance Wu"
618                         elif predictPCA == 7:
619                             name = "Jonathan Pryce"
620                         elif predictPCA == 8:
621                             name = "Max Irons"
622                         elif predictPCA == 9:
623                             name = "Gleen Close"
624                     else:
625                         name = 'unknown'
626                     cv2.putText(img, str(name), (x, y + h), font, 0.55, (0, 255, 0), 1)
627                     i += 1
628                     cv2.imshow('img',img)
629                     key = cv2.waitKey(5)
630                     if key == ord('q'):

```

```
631     |         |     break
632     |         | # if cv2.waitKey(100) & 0xFF == ord('q'):
633     |         |     #     break
634     |         |     video.release()
635     |         |     cv2.destroyAllWindows()
636
637
638
639
640 if __name__ == '__main__':
641     h = haarcascade()
642     h.testshot()
643     # d = dlib()
644     # d.videotest()
```